

Bisync.cloud — Customer AI cost: light → heavy

Planning estimate · 2026-08-05 · Per Bisync company (tenant) · Daily & monthly · Amazon Bedrock AgentCore calculator · USD On-Demand · Excludes tax, Support, shared KB

~\$5	~\$20	~\$61	~\$137
Light \$/tenant / month	Moderate \$/tenant / month	Heavy \$/tenant / month	Very heavy \$/tenant / month

Workloads linked to each customer's Bisync data: **current revenue** Q&A; **purchase/sales forecast** (internal + external); **vendor→component** tags; **reports & how-to**. Quote band: **~\$5 to ~\$140 / tenant / month**.

1. Daily usage per tenant

Workload	Light	Moderate	Heavy	Very heavy
A Revenue Q&A;	3	12	35	80
B Forecast	0	1	4	10
C Tag suggest	2	8	25	50
D Reports / how-to	4	12	30	60
C' Batch SKUs / day	5	30	120	400
Sessions / day	~9.5	~36	~106	~240
Sessions / month	~285	~1,080	~3,180	~7,200

Light = quiet single site · Moderate = small multi-site · Heavy = busy group · Very heavy = HQ + many outlets.

2. Cost per tenant (tokens + AgentCore)

Scenario	Sess/day	\$/day	\$/month	Model \$/mo	AgentCore \$/mo
Light	~9.5	~\$0.16	~\$4.90	~\$4.40	~\$0.45
Moderate	~36	~\$0.68	~\$20	~\$18	~\$2.00
Heavy	~106	~\$2.02	~\$61	~\$54	~\$6.40
Very heavy	~240	~\$4.58	~\$137	~\$123	~\$15

3. Heavy scenario — by use case (/month)

Use case	Sessions/day	\$/month	Share
A Current revenue	35	~\$33	~54%
B Forecast (+ Web Search)	4	~\$9	~15%
C Tags (interactive + batch)	25+12	~\$4	~7%
D Reports / how-to	30	~\$10	~16%
AgentCore overhead	—	~\$6	~10%
Total Heavy tenant	~106	~\$61	100%

4. Fleet rollups (monthly)

Mix	Composition	~\$/month
Pilot	20 × Light	~\$100
Early	50 × Moderate	~\$1,000
Growth	100 × Moderate + 20 × Heavy	~\$3,200
Scale	200 × Moderate + 50 × Heavy + 10 × Very heavy	~\$8,400
Aggressive	300 × Heavy	~\$18,300

5. Suggested commercial AI add-on

Plan	Maps to	AWS cost / mo	Suggested list
Starter	Light	~\$5	\$19 / mo
Standard	Moderate	~\$20	\$49 / mo
Plus	Heavy	~\$61	\$99 / mo
Enterprise	Very heavy (fair use)	~\$137	\$199 / mo

6. AgentCore calculator inputs (per tenant / month)

Component	Light	Moderate	Heavy	Very heavy
Runtime sessions	~285	~1,080	~3,180	~7,200
Gateway invokes	~710	~2,700	~8,000	~18,000
STM events	~855	~3,200	~9,500	~21,600
Web Search queries	0	~45	~180	~450
Code Interpreter	0	~30	~120	~300

Enter one scenario in `calculator.aws/#!/createCalculator/bedrockagentcore`, then × tenant count. Add Bedrock tokens separately (Sonnet for A/B; Haiku for C/D). Shared OpenSearch/KB (~\$400–830/env) is not allocated per tenant.

7. Rates used

Sonnet 4.6 \$3/\$15 per 1M in/out · Haiku 4.5 \$1/\$5 per 1M · Runtime \$0.0895/vCPU-hr + \$0.00945/GB-hr · Web Search \$7/1k · Gateway \$0.005/1k · Memory STM \$0.25/1k · ~20–30% prompt cache on repeated context · ~16 s active CPU / session @ 1 vCPU · peak ~1.5 GB.

Disclaimer: AWS Pricing Calculator estimates exclude tax. Actual fees depend on real sessions, token lengths, cache hit rate, and account terms. Detail markdown: `docs/Bisync-AgentCore-AI-Light-to-Heavy.md`