

Bisync.cloud — Customer AI on Amazon Bedrock AgentCore

Planning estimate · Generated 2026-08-05 · Region baseline ap-southeast-1 · Calculator: [calculator.aws/#/createCalculator/bedrockagentcore](#) · Builds on prior 1,000-location AWS architecture estimate

~\$2,500	~\$3,000	~\$36k	~\$200
Prod AI / mo (1,000 locations)	UAT+Prod AI / mo (UAT ≈ 20% traffic)	UAT+Prod AI / yr	AgentCore platform (ex tokens) / Prod / mo

Prior PDF AI add-on was ~\$2,268 / env / mo (Bedrock + OpenSearch). AgentCore for the customer workloads below stays in the same band; **model tokens + knowledge store dominate**; AgentCore Runtime/Gateway/Memory are comparatively small.

1. Customer-linked workloads

Usage is tied to each customer’s Bisync data: current revenue Q&A; purchase/sales forecast from internal + external factors; vendor-product → component suggestive tags by trend; and report / how-to guidance.

Need	Agent pattern	Model
A Current revenue questions	Tool-using Q&A; over live revenue/sales APIs	Claude Sonnet
B Purchase / sales forecast	Internal tools + Code Interpreter + Web Search	Claude Sonnet
C Vendor → component tags	Trend-aware suggestive tagging (+ batch jobs)	Claude Haiku
D Reports & how-to	Guided help + NL report assembly	Haiku 70% / Sonnet 30%

2. Steady-state volume (Prod, 1,000 locations)

Workload	Sessions/mo	Tokens in/out	Tools / session
A Revenue Q&A;	25,000	6,000 / 1,200	Gateway ×2–3
B Forecast	8,000	10,000 / 2,500	GW ×3 + Code + Web ×1–2
C Tag suggest	15,000	2,500 / 400	Gateway ×1–2
D How-to / reports	20,000	4,000 / 800	Gateway ×1 + KB
C’ Batch tagging	5,000 (~50k SKUs)	1,500 / 200	Gateway ×1
Total sessions	~73,000	—	Prior PDF used ~58k

3. AgentCore calculator line items (Prod)

Component	Calculator input	~/mo USD
Runtime	73k sessions · ~16 s active @ 1 vCPU · 1.5 GB × ~40 s	~\$41
Gateway	~183k InvokeTool @ \$0.005 / 1k	~\$1
Memory STM	~219k events @ \$0.25 / 1k	~\$55
Memory LTM	5k records + ~29k retrievals	~\$19
Code Interpreter	8k forecast sessions	~\$5
Web Search	12k queries @ \$7 / 1k	~\$84
Policy	~183k auth @ \$0.000025	~\$5
Identity	Via Runtime/Gateway	\$0
Browser / Payments	Not used	\$0
AgentCore subtotal	Platform only (ex model tokens)	~\$210

4. Cost roll-up

Line	~/mo USD
Bedrock tokens (Sonnet + Haiku, ~20–30% prompt cache)	~\$1,690
AgentCore platform (Runtime, Memory, Gateway, Policy, Code, Web Search)	~\$210
Knowledge store (OpenSearch Serverless or Bedrock KB)	~\$400–830
CloudWatch observability (AI traces)	~\$50
Prod AI total (mid)	~\$2,350–2,780 → plan ~\$2,500
UAT AI @ ~20% of Prod	~\$500
UAT + Prod AI monthly / yearly	~\$3,000 / ~\$36,000

5. Cost by customer use case (Prod mid)

Use case	Primary drivers	~/mo
A Current revenue	Sonnet tokens + Gateway	~\$780
B Forecast	Sonnet + Web Search + Code Interpreter	~\$590
C Vendor→component tags	Haiku interactive + batch	~\$190
D Reports / how-to	Haiku/Sonnet mix + shared KB	~\$220–400
Shared AgentCore overhead	Runtime / Memory / Policy	~\$200

6. Scenario bands

Scenario	Sessions/mo	Prod AI \$/mo
Pilot (~100 locations)	~7–8k	~\$250–350
Steady 1,000 locations	~73k	~\$2,500
Stretch ~5,000 locations	~250–300k	~\$6.5k–9k

7. Calculator steps

- Open [calculator.aws/#!/createCalculator/bedrockagentcore](#).
- Runtime: 73,000 sessions; ~16 s active CPU @ 1 vCPU; 1.5 GB memory × ~40 s.
- Gateway: ~183,000 API invocations / month.
- Memory: ~219k STM events; 5k LTM records; ~29k retrievals.
- Code Interpreter: 8,000 sessions (forecast only).
- Web Search: 12,000 queries.
- Policy (optional): ~183,000 authorization requests.
- Add Bedrock foundation-model tokens separately (Sonnet/Haiku volumes in §2) — tokens are billed outside AgentCore.
- Keep OpenSearch / Knowledge Base from the prior 1,000-location estimate for how-to RAG.

8. Cost levers

Route how-to + tagging to Haiku; reserve Sonnet for revenue narrative and forecasts. Cache system prompts and report schemas. Cap Web Search (largest AgentCore line after tokens). Batch-tag offline with Haiku. TTL long-term memory per company. Sample AgentCore Evaluations ≤5% if enabled later.

Disclaimer: AWS Pricing Calculator provides an estimate of AWS fees and does not include taxes. Actual fees depend on real usage, region availability, token lengths, cache hit rate, Knowledge Base choice, and account terms. Rates as of AgentCore public pricing (Runtime \$0.0895/vCPU-hr, \$0.00945/GB-hr; Web Search \$7/1k; Gateway \$0.005/1k invocations; Memory STM \$0.25/1k events).

Prior infra+AI share link: [calculator.aws/#!/estimate?id=ed93e08b12ece30328c60d0d20a6a165baf25952](#) · Detail: [docs/Bisync-AgentCore-AI-Usage-Estimate.md](#)